

WAPS: Weighted and Projected Sampling ^{***}

Rahul Gupta¹, Shubham Sharma¹, Subhajit Roy¹, and Kuldeep S. Meel²

¹ Indian Institute of Technology Kanpur, India
`{grahul, smsharma, subhajit}@iitk.ac.in`

² National University of Singapore, Singapore
`meel@comp.nus.edu.sg`

Abstract. Given a set of constraints F and a user-defined weight function W on the assignment space, the problem of constrained sampling is to sample satisfying assignments of F conditioned on W . Constrained sampling is a fundamental problem with applications in probabilistic reasoning, synthesis, software and hardware testing. Consequently, the problem of sampling has been subject to intense theoretical and practical investigations over the years. Despite such intense investigations, there still remains a gap between theory and practice. In particular, there has been significant progress in the development of sampling techniques when W is a uniform distribution, but such techniques fail to handle general weight functions W . Furthermore, we are, often, interested in Σ_1^1 formulas, i.e., $G(X) := \exists Y F(X, Y)$ for some F ; typically the set of variables Y are introduced as auxiliary variables during encoding of constraints to F . In this context, one wonders *whether it is possible to design sampling techniques whose runtime performance is agnostic to the underlying weight distribution and can handle Σ_1^1 formulas?*

The primary contribution of this work is a novel technique, called WAPS, for sampling over Σ_1^1 whose runtime is agnostic to W . WAPS is based on our recently discovered connection between knowledge compilation and uniform sampling. WAPS proceeds by compiling F into a well studied compiled form, d-DNNF, which allows sampling operations to be conducted in linear time in the size of the compiled form. We demonstrate that WAPS can significantly outperform existing state-of-the-art weighted and projected sampler **WeightGen**, by up to 3 orders of magnitude in runtime while achieving a geometric speedup of $296\times$ and solving 564 more instances out of 773. The distribution generated by WAPS is statistically indistinguishable from that generated by an ideal weighted and projected sampler. Furthermore, WAPS is almost oblivious to the number of samples requested.

* Open source tool is available at <https://github.com/meelgroup/waps>.

** This work was supported in part by NUS ODPRT Grant R-252-000-685-133 and AI Singapore Grant R-252-000-A16-490. The computational work for this article was performed on resources of the National Supercomputing Centre, Singapore <https://www.nscc.sg>.

1 Introduction

Boolean satisfiability (SAT) has gathered applications in bounded model checking of hardware and software systems [5,7,51], classical planning [35] and scheduling [27]. Despite the worst-case hardness of SAT, the past few decades have witnessed a significant improvement in the runtime performance of the state-of-the-art SAT solvers [41]. This improvement has led to the usage of SAT solvers as oracles to handle problems whose complexity lies beyond NP. Among these problems, *constrained sampling*, that concerns with sampling from the space of solutions of a set of constraints F , subject to a user-defined weight function W , has witnessed a surge of interest owing to the wide range of applications ranging from machine learning, probabilistic reasoning, software and hardware verification to statistical physics [3,32,39,45].

Not surprisingly, the problem of sampling is known to be computationally intractable. When the weight function W is fixed to a uniform distribution, the problem of constrained sampling is also known as uniform sampling. Uniform sampling has witnessed a long-standing interest from theoreticians and practitioners alike [4,33,38,45]. The past few years, however, have witnessed a significant improvement in the runtime performance of the sampling tools when the weight function W is fixed to a uniform distribution owing to the rise of hashing-based paradigm [2,11,13,22]. While the significant progress for uniform sampling has paved the way for its usage in constrained random simulation [45], the restriction of uniform distribution is limiting, and several applications of constrained sampling require the underlying techniques to be able to handle a wide variety of distributions and related problem formulations as listed below:

Literal-Weighted Sampling In case of literal-weighted sampling, we consider the weight function over assignments defined as the product of the weight of literals, which is specified using a weight function $W(\cdot)$ that assigns a non-negative weight to each literal l in a boolean formula F . As argued in [12], literal-weighted weight function suffices for most of the practical applications ranging from constrained random simulation, probabilistic reasoning, and reliability of power-grids [10,14,21,45].

Projected Sampling Typically, users define constraints in high-level modeling languages such as Verilog [1], Bayesian networks [14] and configuration of grids [21] and then CNF encodings are employed to convert them into a CNF [6]. Commonly used schemes like *Tseitin encoding* [50] introduce auxiliary variables during encoding; though the encoded formulas are equisatisfiable, they typically do not preserve the number of satisfying assignments. In particular, given an initial set of constraints G expressed over a set of variables X , we obtain another formula F such that $G(X) = \exists Y F(X, Y)$. Therefore, we are concerned with sampling over solutions of F *projected over a subset of variables* (such as X in this case). In other words, we are concerned with sampling over Σ_1^1 formulas.

Conditioned Sampling Given a boolean formula Φ and a partial assignment σ , conditioned sampling refers to sampling from the models of Φ that satisfy σ . Conditioning has interesting applications in testing where one is interested in fuzzing the system with inputs that satisfy certain patterns (preconditions). Conditioning has been applied in the past for fault diagnosis [23], conformant planning [46] and databases [15].

Typically, practical applications require sampling techniques that can handle all the above formulations. While techniques based on interval propagation, binary decision diagrams and random perturbation of solution space [22,25,44] cannot handle projection, conditioned, and weighted sampling efficiently, the hashing-based techniques have significantly improved the scalability of sampling techniques and are capable of handling projection and literal-weighted scheme [11,42]. However, the performance of hashing-based techniques is extremely limited in their ability to handle literal-weighted sampling, and one observes a drastic drop in their performance as the weight distribution shifts away from uniform. In this context, one wonders: *whether it is possible to design techniques which can handle projection, conditioned, and literal-weighted sampling without degradation in their performance?*

In this work, we answer the above question in affirmative: we extend our previously proposed knowledge compilation framework in the context of uniform sampling to handle all the three variants. We have implemented a prototype of our framework, named WAPS, and demonstrate that within a time limit of 1800 secs, WAPS performs better than the current state-of-the-art weighted and projected sampler WeightGen [10], by up to 3 orders of magnitude in terms of runtime while achieving a geometric speedup of $296\times$. Out of the 773 benchmarks available, WAPS was able to sample from 588 benchmarks while WeightGen was able to sample from only 24 benchmarks. Furthermore, WAPS is almost oblivious to the number of samples requested.

A significant advantage of our framework is its simplicity: we show that our previously proposed framework in the context of uniform sampling, KUS [49], can be lifted to handle literal-weighted, projection and conditioned sampling. We demonstrate that unlike hashing-based techniques, the runtime performance of WAPS is not dependent on the underlying weight distribution. We want to assert that the simplicity of our framework, combined with its runtime performance and its ability to be agnostic to the underlying distribution is a significant novel contribution to the area of constrained sampling. Besides, an important contribution of our work is the theoretical analysis of sampling techniques that employ knowledge compilation.

The rest of the paper is organized as follows. We first discuss the related work in section 2. We then introduce notations and preliminaries in section 3. In section 4 we present WAPS and do theoretical analysis of WAPS in section 5. We then describe the experimental methodology and discuss results in section 6. Finally, we conclude in section 7.

2 Related Work

Weighted sampling is extensively studied in the literature with the objective of providing scalability while ensuring strong theoretical guarantees. Markov Chain Monte Carlo (MCMC) sampling [32,40] is the most popular technique for weighted sampling; several algorithms like Metropolis-Hastings and simulated annealing have been extensively studied in the literature [36,40]. While MCMC based sampling is guaranteed to converge to a target distribution under mild requirements, convergence is often impractically slow [31]. The practical adaptations for MCMC-based sampling in the context of constrained-random verification has been proposed in [37]. Unfortunately, practical MCMC based sampling tools use heuristics that destroy the theoretical guarantees. Interval-propagation and belief networks have also been employed for sampling [20,26,29], but, though these techniques are scalable, the generated distributions can deviate significantly from the uniform distribution, as shown in [38].

To bridge the wide gap between scalable algorithms and those that give strong guarantees of uniformity several hashing-based techniques have been proposed [10,11,24,28] for weighted sampling. The key idea behind hashing-based techniques is to employ random parity constraints as pairwise independent hash functions to partition the set of satisfying assignments of CNF formula into cells. The hashing-based techniques have achieved significant runtime performance improvement in case of uniform sampling but their scalability suffers for weight distribution and depends strongly on parameters such as *tilt*, which are unlikely to be small for most practical distributions [42].

In recent past, a significant amount of work has been done to compile propositional theory, often represented as a propositional formula in CNF into tractable knowledge representations. One of the prominent and earliest representations is Ordered Binary Decision Diagrams (OBDDs), which have been effectively used for circuit analysis and synthesis [9]. Another family of representations known as Deterministic Decomposable Negation Normal Form (d-DNNF) [19] have proved to be influential in many probabilistic reasoning applications [14,17,18]. Recently, another representation called as Sentential Decision Diagram (SDD) [16] was proposed which maintains canonicity and polytime support for boolean combinations and bridged the gap of succinctness between OBDDs and d-DNNFs. In our recent work [49], we were able to tackle the problem of uniform sampling by exploiting the properties of d-DNNF. Specifically, we were able to take advantage of recent advancements made in the field of knowledge compilation and use the compiled structure to generate uniform samples while competing with the state-of-the-art tools for uniform sampling.

3 Notations and Preliminaries

A literal is a boolean variable or its negation. A clause is a disjunction of a set of literals. A propositional formula F in conjunctive normal form (CNF) is a conjunction of clauses. Let $Vars(F)$ be the set of variables appearing in F . The

set $\text{Vars}(F)$ is called *support* of F . A *satisfying assignment* or *witness* of F , denoted by σ , is an assignment of truth values to variables in its support such that F evaluates to true. We denote the set of all witnesses of F as R_F . Let $\text{var}(l)$ denote the variable of literal l , i.e., $\text{var}(l) = \text{var}(\neg l)$ and $F|_l$ denotes the formula obtained when literal l is set to true in F . Given an assignment σ over $\text{Vars}(F)$ and a set of variables $P \subseteq \text{Vars}(F)$, define $\sigma_P = \{l \mid l \in \sigma, \text{var}(l) \in P\}$ and $R_{F \downarrow P}$ to be the projection of R_F onto P , i.e., $R_{F \downarrow P} = \{\sigma_P \mid \sigma \in R_F\}$.

Given a propositional formula F and a weight function $W(\cdot)$ that assigns a non-negative weight to every literal, the weight of assignment σ denoted as $W(\sigma)$ is the product of weights of all the literals appearing in σ , i.e., $W(\sigma) = \prod_{l \in \sigma} W(l)$. The weight of a set of assignments Y is given by $W(Y) = \sum_{\sigma \in Y} W(\sigma)$. Note that, we have overloaded the definition of weight function $W(\cdot)$ to support different arguments – a literal, an assignment and a set of assignments. We want to highlight that the assumption about weight distribution being generated solely by a literal-weighted function stands well, as many real-world applications like probabilistic inference can be efficiently reduced to literal-weighted sampling [14]. Also, for notational convenience, whenever the formula F , weight function W and sampling set P is clear from the context, we omit mentioning it.

3.1 Weighted and Projected Generators

A *weighted and projected probabilistic generator* is a probabilistic algorithm that generates a witness from $R_{F \downarrow P}$ with respect to weight distribution generated by weight function W . A *weighted and projected generator* $\mathcal{G}^{wp}(\cdot, \cdot, \cdot)$ is a probabilistic generator that guarantees

$$\forall y \in R_{F \downarrow P}, \Pr[\mathcal{G}^{wp}(F, P, W) = y] = \frac{W(y)}{W(R_{F \downarrow P})},$$

An *almost weighted and projected generator* $\mathcal{G}^{awp}(\cdot, \cdot, \cdot)$ relaxes this requirement, ensuring that: given a tolerance $\varepsilon > 0$, $\forall y \in R_{F \downarrow P}$ we have

$$\frac{W(y)}{(1 + \varepsilon)W(R_{F \downarrow P})} \leq \Pr[\mathcal{G}^{awp}(F, P, W) = y] \leq \frac{(1 + \varepsilon)W(y)}{W(R_{F \downarrow P})},$$

Probabilistic generators are allowed to occasionally “fail” in the sense that no witness may be returned even if $R_{F \downarrow P}$ is non-empty. The failure probability for such generators must be bounded by a constant strictly less than 1.




Fig. 1: Example of d-DNNF




Fig. 2: The projected d-DNNF of ex. 1

3.2 Deterministic Decomposable Negation Normal Form (d-DNNF)

To formally define d-DNNF, we first define the Negation Normal Form (NNF):

Definition 1. [19] *Let X be the set of propositional variables. A sentence in NNF is a rooted, directed acyclic graph (DAG) where each leaf node i is labeled with true, false, x or $\neg x$, $x \in X$; and each internal node is labeled with \vee or \wedge and can have arbitrarily many children.*

d-DNNF further imposes that the representation is:

- **Deterministic:** An NNF is deterministic if the operands of \vee in all well-formed boolean formula in the NNF are mutually inconsistent.
- **Decomposable:** An NNF is decomposable if the operands of \wedge in all well-formed boolean formula in the NNF are expressed on a mutually disjoint set of variables.

The deterministic and decomposable properties are conveniently expressed by AND-OR graphs (DAGs) where a node is either an AND node, an OR node or a literal. The operands of AND/OR nodes appear as children of the node. Figure 1 shows an example of d-DNNF representation. For every node t , the subformula corresponding to t is the formula corresponding to d-DNNF obtained by removing all the nodes u such that there does not exist a path from t to u . $T(t)$ represents the set of all partial satisfying assignments for the subformula corresponding to t . The siblings of a node t are the children of the parent of t excluding t itself and the set of such children is given by $Siblings(t)$.

Decision-DNNF is a subset of d-DNNF where the deterministic OR nodes are decision nodes [18]. The state-of-the-art d-DNNF construction tools like C2D [18], DSHARP [43] and D4 [30], construct the Decision-DNNF representation where each OR node has exactly two children while an AND node may have multiple children. Since our framework WAPS employs modern d-DNNF compilers, we assume that the OR node has exactly two children. This assumption is only for the simplicity of exposition as our algorithms can be trivially adopted to the general d-DNNF representations.

4 Algorithm

In this section, we discuss our primary technical contribution: WAPS, weighted and projected sampler that samples from $R_{F \downarrow P}$ with respect to weight function W by employing the knowledge compilation techniques.

WAPS takes a CNF formula F , a set of sampling variables P , a function assigning weights to literals W and required number of samples s as inputs and returns **SampleList**, a list of size s which contain samples such that each sample is independently drawn from the weighted distribution generated by W over $R_{F \downarrow P}$.

Similar to KUS, WAPS (Algorithm 1) mainly comprises of three phases: Compilation, Annotation and Sampling. For d-DNNF compilation, WAPS invokes a specialized compilation routine **PCompile** over the formula F and the sampling

set P (line 1). This is followed by the normalization of weights such that for any literal l , $W'(l) + W'(\neg l) = 1$, where W' is the normalized weight returned in line 2. Then for annotation, **WAnnotate** is invoked in line 3 which uses the weight function W' to annotate weights to all the nodes of the d-DNNF tree. Finally, subroutine **Sampler** (line 4) is invoked which returns s independently drawn samples over P following the weighted distribution generated by W over $R_{F \downarrow P}$. We now describe these three phases in detail.

4.1 Compilation

The compilation phase is performed using the subroutine **PCompile**. **PCompile** is a modified procedure over the component caching and clause learning based algorithm of the d-DNNF compiler **DSHARP** [43,47]. It is presented in Algorithm 2. The changes from the existing algorithm are underlined. The rest of the procedure which is similar to **DSHARP** is mentioned here for completeness. The description of **PCompile** is as follows:

PCompile takes in a CNF formula F in the clausal form and a set of sampling variables P as input and returns a d-DNNF over P . If the formula does not contain any variable from P , **PCompile** invokes **SAT** (line 2) which returns a *True* node if the formula is satisfiable, else it returns a *False* node. Otherwise, **DecideLiteral** is invoked to choose a literal appearing in F such that $\text{var}(l) \in P$ (line 3). This decision is then recursively propagated by invoking **CompileBranch** to create t_1 , the d-DNNF of $F|_l$ (line 4) and t_2 , the d-DNNF of $F|_{\neg l}$ (line 5). **Disjoin** is invoked in line 6 which takes t_1 and t_2 as input and returns t_2 if t_1 is *False* node, t_1 if t_2 is *False* node otherwise a new tree composed by an OR node as the parent of t_1 and t_2 . The result of **Disjoin** is then stored in the cache (line 7) and returned as an output of **PCompile** (line 8).

We now discuss the subroutine **CompileBranch**. It is presented in Algorithm 3. It takes in a CNF formula F , set of sampling variables P and literal l as input and returns a d-DNNF tree of $F|_l$ on P . It first invokes **BCP** (Binary Constraint Propagation), with F , l and P which performs unit-propagation to return a tuple of reduced formula F' and a set of implied literals (*term*) projected over variables in P (line 1). Then **CompileBranch** checks if F' contains an empty clause and returns *False* node to indicate that $F|_l$ is not satisfiable, else the formula is solved using component decomposition as described below.

At line 6 it breaks the formula F' into separate components formed by disjoint set of clauses such that no two components share any variables. Then each component is solved independently (lines 8–15). For each component, it first examines the cache to see if this component has been solved earlier and is present in the cache (line 9). If cache lookup fails, it solves the component with a recursive call to **PCompile** (line 11). If any component is found to be unsatisfiable, *False* node is returned implying that the overall formula is unsatisfiable too, else **CompileBranch** simply conjoins the components' d-DNNFs together with the decided l and implied literals (*term*) and returns this after storing it in the cache for the formula $F|_l$ (lines 16–18).

We illustrate **PCompile** procedure on the following example formula F :

Example 1. $F = \{\{x_1, x_2\}, \{\neg x_3, \neg x_5, x_6\}, \{\neg x_2, x_4, \neg x_1\}, \{x_3, \neg x_6, \neg x_1\}, \{x_6, x_5, \neg x_1, x_3\}, \{x_3, x_6, \neg x_5, \neg x_1\}\}$

Figure 2 represents the d-DNNF of example 1 on $P = \{x_1, x_2, x_3\}$.

Let $P = \{x_1, x_2, x_3\}$ be the set of sampling variables and $Vars(F) = \{x_1, x_2, x_3, x_4, x_5, x_6\}$ be the set of all variables in the formula F . We represent F in clausal form, following steps are then performed for compilation of formula F to a d-DNNF on P :

1. We decide on x_1 which splits the problem into two branches. $C_1 = \{\{x_2\}, \{\neg x_3, \neg x_5, x_6\}\}$, $C_2 = \{\{\neg x_3, \neg x_5, x_6\}, \{\neg x_2, x_4\}, \{x_3, \neg x_6\}, \{x_6, x_5, x_3\}, \{x_3, x_6, \neg x_5\}\}$ as per the positive and negative assignment to x_1 respectively.
2. Then, while solving C_1 , x_2 gets implied by unit propagation and we represent it as a child of AND node. Then we make a decision on x_3 to split the formula into two components. Since both of them turn out to be satisfiable, it is implied that both the assignments of x_3 are possible in this configuration. Hence, we don't have to construct a node for x_3 .
3. Now, to solve C_2 , we make decision on x_3 and find that the component generated on giving negative assignment to x_3 is unsatisfiable. Thus implying positive assignment for x_3 which is represented by child of AND node. This results in the residual formula $C_{2,2} = \{\{\neg x_2, x_4\}, \{\neg x_5, x_6\}\}$
4. For solving $C_{2,2}$, we decide on x_2 to split the formula into two components that turn out to be satisfiable, thus implying that both the assignments of x_2 are possible in this configuration. Hence, we don't have to construct a node for x_2 .

Algorithm 1 WAPS(F, P, W, s)

```

1: dag  $\leftarrow$  PCompile( $F, P$ )
2:  $W' \leftarrow$  Normalize( $W$ )
3: WAnnotate(dag.root,  $W'$ )
4: SampleList  $\leftarrow$  Sampler(dag.root,  $s$ )
5: return SampleList

```

Algorithm 2 PCompile(F, P)

```

1: if  $Vars(F) \cap P = \emptyset$  then
2:   return SAT( $F$ )
3:  $l \leftarrow$  DecideLiteral( $F, P$ )
4:  $t_1 \leftarrow$  CompileBranch( $F, P, l$ )
5:  $t_2 \leftarrow$  CompileBranch( $F, P, \neg l$ )
6:  $t \leftarrow$  Disjoin( $t_1, t_2$ )
7: CacheStore( $F, t$ )
8: return  $t$ 

```

Algorithm 3 CompileBranch(F, P, l)

```

1: ( $F', term$ )  $\leftarrow$  BCP( $F, l, P$ )
2: if  $\emptyset \in F'$  then
3:   CacheStore( $F|_l, False$ )
4:   return  $False$   $\triangleright$  CDCL is done
5: else
6:   Comps  $\leftarrow$  DisjointComponents( $F'$ )
7:   dcomps  $\leftarrow$  {}
8:   for  $C \leftarrow$  Comps do
9:     ddnnf  $\leftarrow$  GetCache( $C$ )
10:    if ddnnf is not found then
11:      ddnnf  $\leftarrow$  PCompile( $C, P$ )
12:    dcomps.Add(ddnnf)
13:    if ddnnf =  $False$  then
14:      CacheStore( $F|_l, False$ )
15:      return  $False$ 
16:  $t =$  Conjoin( $l, term, dcomps$ )
17: CacheStore( $F|_l, t$ )
18: return  $t$ 

```

Algorithm 4 Bottom-Up Pass to annotate d-DNNF with weights on literals

```

1: function WAnnotate( $t, W$ )
2:   if  $\text{label}(t) = \text{Literal}$  then
3:      $t.\text{weight} \leftarrow W(t)$ 
4:   else if  $\text{label}(t) = \text{OR}$  then
5:      $t.\text{weight} \leftarrow 0$ 
6:     for  $c \in \{t.\text{left}, t.\text{right}\}$  do
7:       WAnnotate( $c$ )
8:        $t.\text{weight} \leftarrow t.\text{weight} + c.\text{weight}$ 
9:   else if  $\text{label}(t) = \text{AND}$  then
10:     $t.\text{weight} \leftarrow 1$ 
11:    for  $c \in \text{Childrens}(t)$  do
12:      WAnnotate( $c$ )
13:       $t.\text{weight} \leftarrow t.\text{weight} \times c.\text{weight}$ 

```

Algorithm 5 Top Down pass for sampling s samples

```

1: function Sampler( $t, s$ )
2:   SampleList  $\leftarrow \emptyset$ 
3:   if  $\text{label}(t) = \text{OR}$  then
4:     if  $\text{Weight}(t) = 0$  then
5:        $p \leftarrow 0$ 
6:     else
7:        $p \leftarrow \frac{\text{Weight}(t.\text{left})}{\text{Weight}(t)}$ 
8:        $b \sim \text{Binomial}(s, p)$ 
9:       SampleList0  $\leftarrow \text{Sampler}(t.\text{left}, b)$ 
10:      SampleList1  $\leftarrow \text{Sampler}(t.\text{right}, s - b)$ 
11:      SampleList  $\leftarrow \text{Shuffle}(\text{SampleList}_0 \cup \text{SampleList}_1)$ 
12:      return SampleList
13:   else if  $\text{label}(t) = \text{AND}$  then
14:     for  $c \in \text{children}(t)$  do
15:       SampleListc  $\leftarrow \text{Sampler}(c, s)$ 
16:       SampleList  $\leftarrow \text{Stitch}(\text{SampleList}, \text{SampleList}_c)$ 
17:     return SampleList
18:   else  $\triangleright \text{label}(t) == \text{Literal}$ 
19:     for  $i = 1$  to  $s$  do
20:       Append(val( $t$ ), SampleList)
21:     return SampleList

```

4.2 Annotation

The subroutine WAnnotate is presented in Algorithm 4. WAnnotate takes in a d-DNNF dag and a weight function W as inputs and returns an annotated d-DNNF dag whose each node t is annotated with a weight, given by the sum of weights of all the partial assignments represented by subtree rooted at t . The weights

subsequently annotated on children of an OR node indicate the probability with which it would be selected in the sampling phase.

WAnnotate performs a bottom up traversal on d-DNNF **dag** in a reverse topological order. For each node in the d-DNNF **dag**, WAnnotate maintains an attribute, weight, as per the label of the node given as follows:

- Literal (lines 2–3)** : The weight of a literal node is taken as the weight of the literal given by the weight function W .
- OR (lines 4–8)** : The weight of OR node is made equal to the sum of weights of both of its children.
- AND (lines 9–13)** : The weight of AND node is made equal to the product of weights of all its children.

4.3 Sampling

Algorithm 5 takes the annotated d-DNNF **dag** and the required number of samples s and returns a list L of s samples conforming to the distribution of their weights as governed by weight function $W(\cdot)$ given to WAnnotate. The subroutine **Sampler** is very similar to the sampling procedure in our previous work [49] except that we take the annotated weight of the node instead of the annotated count in the previous work as the probability for Bernoulli trials. The description of **Sampler** procedure as per Algorithm 5 is given as follows:

- OR (lines 3–12)** : We perform s Bernoulli trials with probability p proportional to the weights of each of the OR node’s children to determine the number of samples to draw from them following the weighted distribution. We then recursively draw samples from children as per the overall result of trials simulated via Binomial distribution and shuffle the drawn samples to ensure that the samples are randomly permuted.
- AND (lines 13–17)** : Samples are recursively drawn from each of the children of AND node and then concatenated with each other using the subroutine **Stitch**(line 16)
- Literal (lines 18–21)** : We sample the literal s times for each sample. Function $val(t)$ gives us the literal for literal node.

4.4 Conditioned Sampling

The usage of samplers in testing environment necessitates sampling from F conditioned on fixing assignment to a subset of variables. The state of the art techniques, such as those based on universal hashing, treat every query as independent and are unable to reuse computation across different queries. In contrast, the compilation to d-DNNF allows WAPS to reuse the same d-DNNF. In particular, for a given conditioning expressed as conjunction of different literals, i.e., $\hat{C} = \bigwedge_i l_i$.

In particular, instead of modifying computationally expensive d-DNNF, we modify the weight function as follows:

$$\hat{W}(l) = \begin{cases} 0, & l \notin \hat{C} \\ W(l) & \text{otherwise} \end{cases}$$

5 Theoretical Analysis

We now present theoretical analysis of WAPS, which consists of two components: correctness of WAPS and analysis of behavior of **Sampler** on the underlying d-DNNF graph. First, we prove that WAPS is an exact weighted and projected sampler in Theorem 1. To this end, we prove the correctness of our projected d-DNNF dag compilation procedure **PCompile** in Lemma 1. In Lemma 2, we show that **WAnnotate** annotates each node of the d-DNNF with weights that represent the weight of assignments represented by subtree rooted at that node. This enables us to sample as per the weight distribution in the sampling phase which is proved in Theorem 1 using Lemmas 1 and 2. Secondly, further probing into the behavior of subroutine **Sampler**, we provide an analysis of the probability of visiting any node in the d-DNNF dag while sampling. For this, we first find a probability of visiting a node by following a particular path in Lemma 3 and then we use this result to prove an upper bound for the general case of visiting a node from all possible paths in Theorem 2. We believe that this analysis will motivate the researchers to find new ways to speed up or device new methods to find exact or approximate sampling techniques over a given compiled representation.

Lemma 1. *Given a formula F and set of sampling variables P , the tree returned by $\text{PCompile}(F, P)$ is a d-DNNF dag which represents the set of satisfying assignments of the formula F projected over the set of sampling variables P .*

Proof. Consider a d-DNNF tree for the formula F in which decisions are either taken from the set of sampling variables or non sampling variables only when no sampling variables are left. This tree can be constructed by following the compilation routine of component caching and clause learning based d-DNNF compiler DSHARP [43]. Now, consider any path in such a d-DNNF. It corresponds to a satisfying assignment for the formula F . Consider its trace in the d-DNNF till decisions are taken on sampling variables. At this point, moving further deep in the tree can only lead to further decisions or assignments on non-sampling variables. Therefore, if we cut the edge connecting the rest of the tree to d-DNNF on non-sampling variables, we would have a partial satisfying assignment on this path. Further, ignoring non-sampling variables encountered upto this point, we would have the projection of the satisfying assignment on sampling variables. Now, observe that this can be done for all the satisfying assignments on this path, essentially projecting them to a unique partial satisfying assignment. Now, we claim that d-DNNF on non-sampling variables would only be present if the residual formula is satisfiable. We know that d-DNNF for a formula F can only be constructed if and only if there exists a satisfying assignment. Therefore,

checking for satisfiability, has same effect as that of cutting the edges to d-DNNF on non-sampling variables. If the residual formula is found as unsatisfiable, then it would simply return as it does in DSHARP without existence of d-DNNF on non-sampling variables.

Lemma 2. *Every node t in the d-DNNF dag returned by WAnnotate is annotated by $W(T(t))$, where $T(t)$ is the set of all the partial assignments corresponding to the subtree rooted at t .*

Proof. The proof proceeds by induction on construction of d-DNNF dag. The base case is trivial as for d-DNNF dag of size 1, i.e. literal, the weighted model count for this partial assignment is its weight. The base cases holds trivially true at all the leaves. Since the internal nodes in d-DNNF are only constituted by AND and OR nodes, by proving that weighted model count holds at each internal node, we effectively prove that we get weighted model count at the root of the d-DNNF. The induction holds at each AND and OR node as explained below:

OR: The property of determinism implies that the OR node accumulates mutually disjoint solutions from its children. Therefore the total weights of the partially satisfying assignments at the OR node can be given by the sum of weights of both of its children.

AND: The property of decomposability implies that all the partial assignments of AND node can be given by the concatenating partial assignments of the children of the AND node. The weight of any partial assignment given by AND node will be the product of weights of the corresponding partial assignments given by each of its children. Therefore, the sum of weights of all such satisfying assignments of the AND node can be given by the product of weighted model count of its children which essentially represents the sum of weights of partial satisfying assignments for each of its children.

Theorem 1. *For a given F , P and s , SampleList is the list of samples generated by WAPS. Let $\text{SampleList}[i]$ indicate the sample at the i^{th} index of the list. Then for each $y \in R_{F \downarrow P}$, $\forall i \in [s]$, we have $\Pr[y = \text{SampleList}[i]] = \frac{W(y)}{W(R_{F \downarrow P})}$*

Proof. Lemma 1 shows that the d-DNNF returned by PCompile represents the set of satisfying assignments of the formula F projected over the set of sampling variables P . Lemma 2 shows that d-DNNF returned by WAnnotate is annotated with $W(T(t))$ where $T(t)$ is a set of all the partial assignments given by subtree rooted at corresponding node t . To complete the proof we focus on Subroutine Sampler. The proof proceeds by induction on construction of d-DNNF dag. The base case is trivial as for dag of height one, i.e. literal, there is only one solution, so the base case holds true for all the leaf nodes. Since all the internal nodes of d-DNNF contains only AND and OR nodes we complete the induction step by proving both cases as follows:

OR: The property of determinism ensures disjoint solutions; hence, it boils down to the problem of selecting a model uniformly at random from two buckets

(children nodes) having weighted model count w_l and w_r . To do the same, Bernoulli trail is performed with probability $p = \frac{w_l}{w_l + w_r}$, and partial assignment for solution is taken from respective bucket. In order to draw s samples, s Bernoulli trials are simulated by Binomial distribution by given weights and randomly permuting the union of the solution lists returned from all the children.

AND: The property of decomposability allows us to construct a sample selected uniformly according to weights by sampling a solution segment from each of its children and appending them.

Lemma 3. *For a given F and P , let $\text{fol}(\rho)$ be the event of following a path ρ , which start at root and ends at node t , then $\Pr[\text{fol}(\rho)] = \frac{W(T(t)) \times c_\rho}{W(R_{F \downarrow P})}$ where c_ρ is the product of weight of all the OR nodes' siblings encountered in the path ρ from root to t and $T(t)$ is the set of all the partial satisfying assignments represented by subtree rooted at t .*

Proof. To argue about the probability of visiting a node, we can argue about the probability of going through a path from root to that node. Without loss of generality we can assume that the root node of d-DNNF is an OR node. Let $\rho = (O_1 \rightarrow A_1 \rightarrow \dots \rightarrow t_{2n})$ be a particular path from root node to t_{2n} where $t_{2n} = A_n$ and $\rho_i = (O_1 \rightarrow A_1 \rightarrow \dots \rightarrow t_i)$. Let $N(\rho)$ be a function that returns the end node of the path ρ . For any node t , let $\text{type}(t)$ be a function which returns the type of node t which can be AND or OR. Now, let $\Pr[\text{fol}(\rho_i)]$ be the probability of the event of following the path ρ_i , then:

$$\Pr[\text{fol}(\rho_i) \mid \text{fol}(\rho_{i-1})] = \begin{cases} 1, & \text{if } \text{type}(N(\rho_{i-1})) = \text{AND} \\ \frac{W(N(\rho_i))}{W(N(\rho_{i-1}))}, & \text{if } \text{type}(N(\rho_{i-1})) = \text{OR} \end{cases}$$

Note that $N(\rho_{2n}) = A_n$, $N(\rho_{2n-1}) = O_n$, $N(\rho_{2n-2}) = A_{n-1}$ and $N(\rho_{2n-3}) = O_{n-1}$. Therefore,

$$\begin{aligned} \Pr[\text{fol}(\rho_{2n}) \mid \text{fol}(\rho_{2n-2})] &= \Pr[\text{fol}(\rho_{2n}) \mid \text{fol}(\rho_{2n-1})] \times \Pr[\text{fol}(\rho_{2n-1}) \mid \text{fol}(\rho_{2n-2})] \\ &= \frac{W(N(\rho_{2n}))}{W(N(\rho_{2n-1}))} \times 1 \end{aligned}$$

$$\begin{aligned} \Pr[\text{fol}(\rho_{2n}) \mid \text{fol}(\rho_{2n-3})] &= \Pr[\text{fol}(\rho_{2n}) \mid \text{fol}(\rho_{2n-2})] \times \Pr[\text{fol}(\rho_{2n-2}) \mid \text{fol}(\rho_{2n-3})] \\ &= \frac{W(N(\rho_{2n}))}{W(N(\rho_{2n-1}))} \times 1 \times \frac{W(N(\rho_{2n-2}))}{W(N(\rho_{2n-3}))} \\ &= \frac{W(A_n)}{W(O_n)} \times 1 \times \frac{W(A_{n-1})}{W(O_{n-1})} \end{aligned}$$

Now,

$$W(A_{n-1}) = W(O_n) \times \prod_{i \in \text{Siblings}(O_n)} W(i)$$

Using the above results,

$$\begin{aligned}\Pr[\text{fol}(\rho_{2n}) \mid \text{fol}(\rho_{2n-3})] &= \frac{W(A_n)}{W(O_n)} \times \frac{W(O_n) \times \prod_{i \in \text{Siblings}(O_n)} W(i)}{W(O_{n-1})} \\ &= W(A_n) \times \frac{\prod_{i \in \text{Siblings}(O_n)} W(i)}{W(O_{n-1})}\end{aligned}$$

Now, applying the above method recursively, we get

$$\Pr[\text{fol}(\rho_{2n}) \mid \text{fol}(\rho_1)] = \frac{W(A_n) \times \prod_{i \in [1, n]} \prod_{j \in \text{Siblings}(O_i)} W(j)}{W(R_{F \downarrow P})}$$

Since, we always visit the root node to fetch any samples, $\Pr[\text{fol}(\rho_1)] = 1$. Therefore,

$$\begin{aligned}\Pr[\text{fol}(\rho_{2n})] &= \frac{W(A_n) \times \prod_{i \in [1, n]} \prod_{j \in \text{Siblings}(O_i)} W(j)}{W(R_{F \downarrow P})} \\ &= \frac{W(T(t)) \times c}{W(R_{F \downarrow P})}\end{aligned}$$

where $c = \prod_{i \in [1, n]} \prod_{j \in \text{Siblings}(O_i)} W(j)$ i.e the product of weight of all the OR nodes' siblings encountered in a path from root to t_{2n} .

Theorem 2. *For a given F and P , let $\text{visit}(t)$ be the event of visiting a node t to fetch one sample as per subroutine **Sampler**, then $\Pr[\text{visit}(t)] \leq \frac{W(\Gamma(t))}{W(R_{F \downarrow P})}$ where $\Gamma(t) = \{\sigma \mid \sigma \in R_{F \downarrow P}, \sigma_{\downarrow \text{Vars}(t)} \in T(t)\}$ and $T(t)$ is a set of all the partial satisfying assignments represented by subtree rooted at t .*

Proof. In lemma 3 we have calculated the probability of visiting a node t by taking a particular path from root to node t . So the probability of visiting a node t will be the sum of probability of visiting t by all possible paths. Let $\mathcal{P} = \{\rho_1, \rho_2, \dots, \rho_m\}$ be the set of all paths from root to node t and $\text{visit}(t)$ be the event of visiting a node t in subroutine **Sampler** then,

$$\Pr[\text{visit}(t)] = \sum_{\rho \in \mathcal{P}} \Pr[\text{visit}(t) \mid \text{fol}(\rho)] \times \Pr[\text{fol}(\rho)] = \sum_{\rho \in \mathcal{P}} 1 \times \Pr[\text{fol}(\rho)]$$

From Lemma 3, $\Pr[\text{visit}(t)] = \sum_{\rho \in \mathcal{P}} \frac{W(T(t)) \times c_\rho}{W(R_{F \downarrow P})}$ where c_ρ is the product of the weight of all the OR nodes' siblings encountered in a path ρ from root to t . For any such path, we call $\{t_\rho^1, t_\rho^2, \dots, t_\rho^n\}$ as the set of all the OR node siblings encountered on the path ρ . Now, let σ_ρ^{ext} be the set of assignments over P represented by path ρ . Therefore,

$$\sigma_\rho^{\text{ext}} = T(t_\rho^1) \times T(t_\rho^2) \cdots \times T(t_\rho^n) \times T(t)$$

where, $T(\cdot)$ are set of assignments and \times is a cross product. Now, any tuple from σ_ρ^{ext} represents a satisfying assignment in the d-DNNF. Therefore, $\sigma_\rho^{ext} \subseteq R_{F \downarrow P}$. Note that, from lemma 2, it follows that weight annotated by **WAnnotate** at t is equal to $W(T(t))$. Therefore,

$$c_\rho = W(T(t_\rho^1)) \times W(T(t_\rho^2)) \cdots \times W(T(t_\rho^n))$$

And, $W(\sigma_\rho^{ext}) = W(T(t)) \times c_\rho$. Notice that, $\sigma_\rho^{ext} \subseteq \Gamma(t)$ as σ_ρ^{ext} represents satisfying extensions of partial assignments contained in $T(t)$ itself. This is true $\forall \rho \in \mathcal{P}$. Therefore as $W(\cdot)$ is an increasing function,

$$\bigcup_{\rho \in \mathcal{P}} \sigma_\rho^{ext} \subseteq \Gamma(t) \implies \sum_{\rho \in \mathcal{P}} W(T(t)) \times c_\rho \leq W(\Gamma(t))$$

Note that, the inequality indeed holds as the intersection of sets of partial assignments represented by t and any other node not lying on the path from root to t may not be ϕ (empty). Therefore,

$$\sum_{\rho \in \mathcal{P}} \frac{W(T(t)) \times c_\rho}{W(R_{F \downarrow P})} \leq \frac{W(\Gamma(t))}{W(R_{F \downarrow P})} \implies \Pr[\text{visit}(t)] \leq \frac{W(\Gamma(t))}{W(R_{F \downarrow P})}$$

6 Evaluation

In order to evaluate the runtime performance and analyze the quality of samples generated by WAPS, we implemented a prototype in Python. For d-DNNF compilation, our prototype makes use of DSHARP [43] when sampling set is available else we use D4 [30]. We would have preferred to use state-of-the-art d-DNNF compiler D4 but owing to its closed source implementation, we could not modify it as per our customized compilation procedure **PCompile**. Therefore, for projected compilation, we have modified DSHARP which has an open-source implementation. We have conducted our experiments on a wide range of publicly available benchmarks. In all, our benchmark suite consisted of 773 benchmarks arising from a wide range of real-world applications. Specifically, we used constraints arising from DQMR networks, bit-blasted versions of SMT-LIB (SMT) benchmarks, and ISCAS89 circuits [8] with parity conditions on randomly chosen subsets of outputs and nextstate variables [34,48]. We assigned random weights to literals wherever weights were not already available in our benchmarks. All our experiments were conducted on a high performance compute cluster whose each node consists of E5-2690 v3 CPU with 24 cores and 96GB of RAM. We utilized single core per instance of benchmark with a timeout of 1800 seconds.

The objective of our evaluation was to answer the following questions:

1. How does WAPS perform in terms of runtime in comparison to **WeightGen**, the current state-of-the-art weighted and projected sampler?
2. How does WAPS perform for *incremental sampling* and scales when asked for different number of samples?

Table 1: Run time (in seconds) for 1000 samples

Benchmark	Vars	Clauses	$ P $	WeightGen	WAPS			Speedup on WeightGen
					Compile	A+S	Total	
s526a.3.2	366	944	24	490.34	15.37	1.96	17.33	28.29
LoginService	11511	41411	36	1203.93	15.02	0.75	15.77	76.34
blockmap.05.02	1738	3452	1738	1140.87	0.04	5.30	5.34	213.65
s526.3.2	365	943	24	417.24	0.06	0.67	0.73	571.56
or-100-5-4-UC-60	200	500	200	1795.52	0.01	0.74	0.74	2426.38
or-50-5-10-UC-40	100	250	100	1292.67	0.01	0.36	0.36	3590.75
blasted_case35	400	1414	46	TO	0.57	1.46	2.03	-
or-100-20-4-UC-50	200	500	200	TO	0.19	2.48	2.67	-

3. How does the distribution of samples generated by WAPS compare with the distribution generated by an ideal weighted and projected sampler?
4. How does WAPS perform for conditioning on arbitrary variables?
5. How does our knowledge compilation based sampling techniques perform in comparison to hashing based sampling techniques for the task of generalizing to arbitrary weight distributions?

Our experiment demonstrated that within a time limit of 1800 secs, WAPS is able to significantly outperform existing state-of-the-art weighted and projected sampler WeightGen, by up to 3 orders of magnitude in terms of runtime while achieving a geometric speedup of $296\times$. Out of the 773 benchmarks available WAPS was able to sample from 588 benchmarks while WeightGen was able to sample from only 24 benchmarks. For *incremental sampling*, WAPS achieves a geometric speedup of 3.69. Also, WAPS is almost oblivious to the number of samples requested. Empirically, the distribution generated by WAPS is statistically indistinguishable from that generated by an ideal weighted and projected sampler. Also, while performing conditioned sampling in WAPS, we incur no extra cost in terms of runtime in most of the cases. Moreover, the performance of our knowledge compilation based sampling technique is found to be oblivious to weight distribution. We present results for only a subset of representative benchmarks here. Detailed data along with the expanded versions of all the tables presented here is available under Runtime Results directory at <https://github.com/meelgroup/waps>

Number of instances solved. We compared the runtime performance of WAPS with WeightGen [10] (state-of-the-art weighted and projected sampler) by generating 1000 samples from each tool with a timeout of 1800 secs. Figure 3 shows the cactus plot for WeightGen and WAPS. We present the number of benchmarks on the x -axis and the time taken on y -axis. A point (x, y) implies that x benchmarks took less than or equal to y seconds to sample. All our runtime statistics for WAPS include the time for the knowledge compilation phase (via D4 or DSHARP). From all the 773 available benchmarks WeightGen was able to sample from only 24 benchmarks while WAPS was able to sample from 588 benchmarks. Table 1 shows the runtimes of some of the benchmarks on the two tools. The columns in the table give the benchmark name, number of variables,

number of clauses, size of sampling set, time taken in seconds by **WeightGen** and **WAPS** divided into time taken by Compilation and A+S: Annotation and Sampling followed by speedup of **WAPS** with respect to **WeightGen**. Table 1 clearly shows that **WAPS** outperforms **WeightGen** by upto 3 orders of magnitude. For all the 24 benchmarks that **WeightGen** was able to solve **WAPS** outperformed **WeightGen** with a geometric speedup of $296\times$.

Incremental Sampling Incremental sampling involves fetching multiple, relatively small-sized samples until the objective (such as desired coverage or violation of property) is achieved. We benefit from pre-compiled knowledge representations in this scenario, as they allow us to perform repeated sampling as per varied distributions. If weights are changed, we simply Annotate the tree again followed by sampling, else, we directly move to the sampling phase, thus saving a significant amount of time by bypassing the compilation phase.

In our experiments, we have evaluated the time taken by **WAPS** for 1000 samples in 10 successive calls with same weights. The results are presented in table 2 for a subset of benchmarks. The first column mentions the benchmark name with the number of variables, clauses and size of sampling set in subsequent columns. The time taken by **WAPS** for first run to fetch 1000 samples is given in the fifth column while the overall time taken for first run together with the subsequent 9 incremental runs is presented in sixth column. The final column shows the average gain in terms of speedup calculated by taking the ratio of time taken by **WAPS** for first run with the average time taken by **WAPS** for subsequent 9 incremental runs thus resulting in a total of 10000 samples. Overall, **WAPS** achieves a geometric speedup of $3.69\times$ on our set of benchmarks.




Fig. 3: Cactus Plot comparing **WeightGen** and **WAPS**.

Benchmark	Vars	Clauses	P	WAPS		Speedup
				1000	10,000	
case110	287	1263	287	1.14	9.28	1.26
or-70-10-10-UC-20	140	350	140	2.75	9.02	6.56
s526.7.4	383	1019	24	60.38	143.16	13.20
or-60-5-2-UC-10	120	300	120	12.10	20.35	16.50
s35932.15.7	17918	44709	1763	69.01	106.65	20.73
case121	291	975	48	35.85	51.41	20.73
s641.15.7	576	1399	54	729.38	916.83	35.01
squaring7	1628	5837	72	321.95	365.13	67.10
LoginService	11511	41411	36	15.89	18.12	64.13
ProjectService	3175	11019	55	184.51	195.25	154.61



Table 2: Runtimes (in sec.) of **WAPS** for incremental sampling

Effect of number of samples To check how **WAPS** scales with different number of samples, we invoked **WAPS** for fetching different number of samples: 1000, 2000, 4000, 8000, 10000 with a timeout of 1800 secs. Table 3 presents the runtime of **WAPS** for different samples on some benchmarks. The first column represents the benchmark name. Second, third and fourth columns represent the number of variables, clauses and size of sampling set. The next five columns represent the time taken by **WAPS** for 1000, 2000, 4000, 8000 and 10000 samples. Table

Table 3: Runtime (in sec.) of WAPS to generate different size samples

Benchmark	Vars	Clauses	$ P $	Sampling Size				
				1000	2000	4000	8000	10000
s1488.7.4	872	2499	14	0.5	0.75	1.29	2.2	2.9
s444.15.7	377	1072	24	0.74	1.29	1.91	3.46	4.12
s526.3.2	365	943	24	0.84	1.03	1.86	3.71	4.22
s820a.3.2	598	1627	23	0.63	1.03	2.04	3.92	4.81
case35	400	1414	46	2.38	3.22	5.31	9.38	11.41
LoginService	11511	41411	36	15.8	16.12	16.68	18.3	18.36
ProjectService	3175	11019	55	184.22	184.99	188.33	191.16	193.92
or-60-20-6-UC-10	120	300	120	1465.34	1458.23	1494.46	1499.67	1488.23

3 clearly demonstrates that WAPS is almost oblivious to the number of samples requested.

**Fig. 4:** Cactus Plot comparing WeightGen and UniGen2.**Fig. 5:** Cactus Plot comparing WAPS and KUS.

Uniform sampling generalized for weighted sampling To explore the trend in performance between uniform and weighted sampling on the dimension of hashing based techniques pitched against our newly proposed sampling techniques based on knowledge compilation, we compared WAPS to KUS in a parallel comparison between WeightGen and UniGen2. Specifically, we ran WAPS for weighted sampling and KUS for uniform sampling without utilizing the sampling set as KUS does not support the sampling set. On the other hand, for hashing based sampling techniques, we compared WeightGen to UniGen2 while using the sampling set. Figure 4 shows the cactus plot for WeightGen and UniGen2 and figure 5 shows a cactus plot for WAPS and KUS. From all the 773 benchmarks, WeightGen was able to sample from only 24 benchmarks while UniGen2 was able to sample from 208 benchmarks. In comparison, WAPS was able to sample from 606 benchmarks while KUS was able to sample from 602 benchmarks. Our experiments demonstrated that the performance of hashing-based techniques is extremely limited in their ability to handle literal-weighted sampling and there

is a drastic drop in their performance as the weight distribution shifts away from uniform. While for our knowledge compilation based sampling techniques we observe that their performance is oblivious to the weight distribution.

Distribution comparison We measure the distribution of WAPS vis-a-vis an *ideal weighted and projected sampler* (IS) and observed that WAPS is statistically indistinguishable from IS. We implemented the *ideal weighted and projected sampler* (IS) (similar to [10]) such that given a CNF formula F , IS first generates all the satisfying assignments, then compute their weights and uses these weights to sample the ideal distribution. We generated a large number N (6×10^6) of samples using WAPS, WeightGen and IS. In each case, the number of times various samples were generated was recorded, giving a distribution of the counts. Figure 6 shows the distributions generated by WAPS, WeightGen and IS for one of benchmark (case110) which has 16,384 solutions. In the figure, each point (x, y) indicates that x distinct models are generated y times. Since the number of samples is much larger than the number of models, each model occurred multiple times in the list of samples. We computed the frequency of generation of individual model and grouped the models that had the same frequency. Figure 6 shows that the distribution generated by WAPS is practically indistinguishable from IS.




Fig. 6: Distribution comparison between IS, WAPS and WeightGen on case110 with 6×10^6 samples.

Effect of conditioning on variables We evaluated the performance of WAPS in the context of conditioned sampling i.e. partial assignments enforced on some

of the CNF’s variables apart from the usual CNF clauses. In particular, we experimented by randomly choosing 1%, 2%, 3%, 4% and 5% of all variables 10 times and assigned arbitrary signs to chosen variables. This resulted in 50 instances for each problem with different conditions and we sampled 1000 samples from each instance. For each instance, let the set of conditions be represented by a set of literals L . We observed an improvement in average runtime as more and more variables get constrained since, the paths which lead to assignments not conforming to the given condition get pruned in the annotation phase. This happens, as some nodes in d-DNNF are attributed with 0 weight. This occurs when all the partial assignments represented by the subtree rooted at a node in d-DNNF becomes inconsistent with the added conditions. Table 4 represent the results of conditioned sampling on representative set of benchmarks. The first column represents the benchmark name. Second, third and fourth columns represent the number of variables, clauses and d-DNNF compilation time. The next five columns represent the time taken by WAPS when the formula from benchmark is conditioned on 1%, 2%, 3%, 4% and 5% of all variables. We report the average time over 10 runs for each case. The acronym NS for $y\%$ conditioned variables means that no samples could be drawn for atleast one such instance as it became unsatisfiable by the additional conditions.

Table 4: Runtime (in sec.) of WAPS to generate samples while conditioning on arbitrary $y\%$ variables

Benchmark	Vars	Clauses	WAPS(Compilation and A+S on $y\%$ conditioned vars)					
			Compile	$y = 1$	$y = 2$	$y = 3$	$y = 4$	$y = 5$
s27_new_15.7	17	43	0.01	0.29	0.27	NS	0.28	NS
or-50-20-2	100	250	6.77	112.32	111.56	109.17	111.01	110.63
or-50-5-1	100	250	8.91	154.21	153.62	154.86	153.47	152.91
or-50-20-3	100	250	13.47	184.78	187.13	185.68	185.59	185.02
or-50-10-6	100	250	15.36	342.58	339.97	341.09	339.52	337.56
or-50-20-7	100	250	23.93	384.90	381.29	383.07	380.76	380.94

7 Conclusion

In this paper, we designed a knowledge compilation-based framework, called WAPS, for literal-weighted, projected and conditional sampling. WAPS provides strong theoretical guarantees and its runtime performance upon the existing state-of-the-art weighted and projected sampler WeightGen, by up to 3 orders of magnitude in terms of runtime. Out of the 773 benchmarks available, WAPS is able to sample from 588 benchmarks while WeightGen is only able to sample from 24 benchmarks. WAPS achieves a geometric speedup of 3.69 for *incremental sampling*. It is worth noting that WeightGen has weaker guarantees than WAPS. Furthermore, WAPS is almost oblivious to the number of samples requested.

References

1. System Verilog. <http://www.systemverilog.org> (2015)
2. Achlioptas, D., Hammoudeh, Z.S., Theodoropoulos, P.: Fast Sampling of Perfectly Uniform Satisfying Assignments. In: Proc. of SAT. pp. 135–147 (2018)
3. Bacchus, F., Dalmao, S., Pitassi, T.: Algorithms and Complexity Results for #SAT and Bayesian Inference. In: Proc. of FOCS. pp. 340–351 (2003)
4. Bellare, M., Goldreich, O., Petrank, E.: Uniform Generation of NP-witnesses using an NP-oracle. *Information and Computation* **163**(2), 510–526 (2000)
5. Biere, A., Cimatti, A., Clarke, E., Fujita, M., Zhu, Y.: Symbolic model checking using SAT procedures instead of BDDs. In: Proc. DAC. pp. 317–320 (1999)
6. Biere, A., Heule, M., van Maaren, H., Walsh, T.: *Handbook of Satisfiability*. IOS Press (2009)
7. Biesse, P., Leonard, T., Mokkedem, A.: Finding bugs in an alpha microprocessor using satisfiability solvers. In: Proc. of CAV. pp. 454–464 (2001)
8. Brglez, F., Bryan, D., Kozminski, K.: Combinational profiles of sequential benchmark circuits. In: Proc. of ISCAS. pp. 1929–1934 (1989)
9. Bryant, R.E.: Symbolic boolean manipulation with ordered binary-decision diagrams. *ACM Computing Surveys (CSUR)* **24**(3), 293–318 (1992)
10. Chakraborty, S., Fremont, D.J., Meel, K.S., Seshia, S.A., Vardi, M.Y.: Distribution-Aware Sampling and Weighted Model Counting for SAT. In: Proc. of AAAI. pp. 1722–1730 (2014)
11. Chakraborty, S., Fremont, D.J., Meel, K.S., Seshia, S.A., Vardi, M.Y.: On Parallel Scalable Uniform SAT Witness Generation. In: Proc. of TACAS. pp. 304–319 (2015)
12. Chakraborty, S., Fried, D., Meel, K.S., Vardi, M.Y.: From weighted to unweighted model counting. In: Proceedings of IJCAI. pp. 689–695 (2015)
13. Chakraborty, S., Meel, K.S., Vardi, M.Y.: A Scalable and Nearly Uniform Generator of SAT Witnesses. In: Proc. of CAV. pp. 608–623 (2013)
14. Chavira, M., Darwiche, A.: On probabilistic inference by weighted model counting. *Artificial Intelligence* **172**(6), 772–799 (2008)
15. Dalvi, N.N., Schnaitter, K., Suciu, D.: Computing query probability with incidence algebras. In: Proc. of PODS. pp. 203–214 (2010)
16. Darwiche, A.: SDD: A new canonical representation of propositional knowledge bases. In: Proc. 22nd Int’l Joint Conf. on Artificial Intelligence. pp. 819–826 (2011)
17. Darwiche, A.: On the tractable counting of theory models and its application to belief revision and truth maintenance. *CoRR* (2000)
18. Darwiche, A.: New Advances in Compiling CNF to Decomposable Negation Normal Form. In: Proc. of ECAI. pp. 318–322 (2004)
19. Darwiche, A., Marquis, P.: A knowledge compilation map. *Journal of Artificial Intelligence Research* **17**, 229–264 (2002)
20. Dechter, R., Kask, K., Bin, E., Emek, R.: Generating random solutions for constraint satisfaction problems. In: Proc. of AAAI. pp. 15–21 (2002)
21. Duenas-Osorio, L., Meel, K.S., Paredes, R., Vardi, M.Y.: Counting-based reliability estimation for power-transmission grids. In: Proc. of AAAI (2017)
22. Dutra, R., Laeuffer, K., Bachrach, J., Sen, K.: Efficient sampling of SAT solutions for testing. In: Proc. of ICSE. pp. 549–559 (2018)
23. Elliott, P., Williams, B.: DNNF-based Belief State Estimation. In: Proc. of AAAI. pp. 36–41 (2006)

24. Ermon, S., Gomes, C.P., Sabharwal, A., Selman, B.: Embed and Project: Discrete Sampling with Universal Hashing. In: Proc. of NIPS. pp. 2085–2093 (2013)
25. Ermon, S., Gomes, C.P., Selman, B.: Uniform Solution Sampling Using a Constraint Solver As an Oracle. In: Proc. of UAI. pp. 255–264 (2012)
26. Gogate, V., Dechter, R.: A New Algorithm for Sampling CSP Solutions Uniformly at Random. In: Proc. of CP. pp. 711–715 (2006)
27. Gomes, C.P., Selman, B., McAloon, K., Tretkoff, C.: Randomization in Backtrack Search: Exploiting Heavy-tailed Profiles for Solving Hard Scheduling Problems. In: Proc. of AIPS (1998)
28. Ivrii, A., Malik, S., Meel, K.S., Vardi, M.Y.: On computing minimal independent support and its applications to sampling and counting. *Constraints* pp. 1–18 (2015)
29. Iyer, M.A.: RACE: A word-level ATPG-based constraints solver system for smart random simulation. In: Proc. of ITC. pp. 299–308 (2003)
30. Jean-Marie Lagniez, P.M.: An Improved Decision-DNNF Compiler. In: Proc. of IJCAI. pp. 667–673 (2017)
31. Jerrum, M.R., Sinclair, A.: Approximating the permanent. *SIAM journal on computing* **18**(6), 1149–1178 (1989)
32. Jerrum, M.R., Sinclair, A.: The Markov Chain Monte Carlo method: an approach to approximate counting and integration. *Approximation algorithms for NP-hard problems* pp. 482–520 (1996)
33. Jerrum, M.R., Valiant, L.G., Vazirani, V.V.: Random generation of combinatorial structures from a uniform distribution. *Theoretical Computer Science* **43**(2-3), 169–188 (1986)
34. John, A.K., Chakraborty, S.: A quantifier elimination algorithm for linear modular equations and disequations. In: Proc. of CAV. pp. 486–503 (2011)
35. Kautz, H., Selman, B.: Pushing the envelope: Planning, propositional logic, and stochastic search. In: Proc. of AAAI (1996)
36. Kirkpatrick, S., Gelatt, C.D., Vecchi, M.P.: Optimization by simulated annealing. *Science* **220**(4598), 671–680 (1983)
37. Kitchen, N.: Markov Chain Monte Carlo Stimulus Generation for Constrained Random Simulation. Ph.D. thesis, University of California, Berkeley (2010)
38. Kitchen, N., Kuehlmann, A.: Stimulus generation for constrained random simulation. In: Proc. of ICCAD. pp. 258–265 (2007)
39. Madras, N., Piccioni, M.: Importance sampling for families of distributions. *Annals of applied probability* pp. 1202–1225 (1999)
40. Madras, N.: Lectures on Monte Carlo Methods, Fields Institute Monographs 16. American Mathematical Society (2002)
41. Malik, S., Zhang, L.: Boolean satisfiability from theoretical hardness to practical success. *Commun. ACM* **52**(8), 76–82 (2009)
42. Meel, K.S.: Constrained Counting and Sampling: Bridging the Gap between Theory and Practice. Ph.D. thesis, Rice University (2017)
43. Muise, C., McIlraith, S.A., Beck, J.C., Hsu, E.: DSHARP: Fast d-DNNF Compilation with sharpSAT. In: Proc. of AAAI. pp. 356–361 (2016)
44. Naveh, R., Metodi, A.: Beyond feasibility: CP usage in constrained-random functional hardware verification. In: Proc. of CP. pp. 823–831 (2013)
45. Naveh, Y., Rimón, M., Jaeger, I., Katz, Y., Vinov, M., Marcu, E., Shurek, G.: Constraint-based random stimuli generation for hardware verification. In: Proc of IAAI. pp. 1720–1727 (2006)
46. Palacios, H., Bonet, B., Darwiche, A., Geffner, H.: Pruning Conformant Plans by Counting Models on Compiled d-DNNF Representations. In: Proc. of ICAPS. pp. 141–150 (2005)

47. Sang, T., Bacchus, F., Beame, P., Kautz, H.A., Pitassi, T.: Combining component caching and clause learning for effective model counting. In: Proc. of SAT (2004)
48. Sang, T., Beame, P., Kautz, H.: Performing Bayesian inference by weighted model counting. In: Proc. of AAAI. pp. 475–481 (2005)
49. Sharma, S., Gupta, R., Roy, S., Meel, K.S.: Knowledge Compilation meets Uniform Sampling. In: Proc. of LPAR-22. pp. 620–636 (2018)
50. Tseitin, G.S.: On the Complexity of Derivation in Propositional Calculus (1983)
51. Velev, M.N., Bryant, R.E.: Effective Use of Boolean Satisfiability Procedures in the Formal Verification of Superscalar and VLIW Microprocessors. *J. Symb. Comput.* (2), 73–106 (2003)